



The Butterfly Effect in artificial intelligence systems: Implications for AI bias and fairness

Emilio Ferrara

University of Southern California, Thomas Lord Department of Computer Science, Los Angeles, 90007, CA, USA

ARTICLE INFO

Keywords:

Artificial intelligence bias
AI fairness
AI ethics

ABSTRACT

The concept of the Butterfly Effect, derived from chaos theory, highlights how seemingly minor changes can lead to significant, unpredictable outcomes in complex systems. This phenomenon is particularly pertinent in the realm of AI fairness and bias. Factors such as subtle biases in initial data, deviations during algorithm training, or shifts in data distribution from training to testing can inadvertently lead to pronounced unfair results. These results often disproportionately impact marginalized groups, reinforcing existing societal inequities. Furthermore, the Butterfly Effect can magnify biases in data or algorithms, intensify feedback loops, and heighten susceptibility to adversarial attacks. Recognizing the complex interplay within AI systems and their societal ramifications, it is imperative to rigorously scrutinize any modifications in algorithms or data inputs for possible unintended effects. This paper proposes a combination of algorithmic and empirical methods to identify, measure, and counteract the Butterfly Effect in AI systems. Our approach underscores the necessity of confronting these challenges to foster equitable outcomes and ensure responsible AI evolution.

1. Introduction

The Butterfly Effect, a fundamental concept in chaos theory, describes the sensitive dependence on initial conditions in nonlinear, dynamic systems. It was coined by American mathematician and meteorologist Edward Lorenz in the early 1960s (Lorenz, 1963). The Butterfly Effect suggests that small initial changes in complex, dynamic systems can result in significantly different and often unpredictable outcomes over time. The idea is best portrayed by the popular saying that the flap of a butterfly's wings in Brazil could set off a chain of events leading to a tornado in Texas. While working on numerical weather prediction models, Lorenz observed that small changes in initial conditions led to drastically different long-term forecasts. The importance of the Butterfly Effect extends beyond meteorology and has found applications in various scientific disciplines, including physics, engineering, biology, and social sciences. In these disciplines, the Butterfly Effect is used to describe the consequences of small changes or perturbations in complex systems: this effect highlights the non-linear nature of complex systems as well as the importance of understanding their interconnectedness of various components and the challenges associated with accurate long-term predictions (Strogatz, 2018).

1.1. Relevance of the Butterfly Effect to AI fairness and bias

In the context of AI fairness and bias (Ferrara, 2023a), the Butterfly Effect highlights the potential for small biases or skewed data inputs

at various stages of algorithm development to result in significant and unexpected unfair outcomes. This phenomenon can manifest in various ways, such as small adjustments to input data, inherent biases within the data or algorithms themselves, shifts in data distributions, adversarial attacks, or feedback loops that amplify existing biases. The interconnected nature of AI systems, combined with their potential impact on society, makes it crucial to examine the Butterfly Effect's role in AI fairness and bias.

1.2. Factors contributing to the Butterfly Effect in AI systems

The relevance of the Butterfly Effect to AI fairness and bias lies in the observation that AI and machine learning (ML) systems are complex and interconnected, with multiple components contributing to their final decision-making process. These components include data, algorithms, and user interactions, among others. As a result, seemingly small changes or biases in these components can have a significant and potentially unforeseen impact on the fairness and bias of AI systems. Furthermore, seemingly small initial biases can propagate and cause large disparities. Vice versa, larger initial biases might possibly have less dramatic consequences in real applications. Several factors can contribute to the emergence of the Butterfly Effect on AI fairness and bias, summarized in Table 1.

E-mail address: emiliofe@usc.edu.

<https://doi.org/10.1016/j.mlwa.2024.100525>

Received 19 September 2023; Received in revised form 6 December 2023; Accepted 2 January 2024

Available online 19 January 2024

2666-8270/© 2024 The Author(s). Published by Elsevier Ltd. This is an open access article under the CC BY license (<http://creativecommons.org/licenses/by/4.0/>).

Table 1
Factors contributing to the Butterfly Effect in AI systems.

Contributing factor	Description	References
High-dimensional input space	ML algorithms often operate on high-dimensional input data, which means they rely on a multitude of features to make decisions. Small perturbations in the input data, such as the removal or addition of features, can lead to vastly different model behavior and predictions. This sensitivity to input data can manifest as a Butterfly Effect, where slight adjustments result in significant and unintended consequences.	Barocas and Selbst (2016)
Nonlinearity and complexity of ML models	Many ML models, such as deep neural networks, are highly nonlinear and complex. This nonlinearity can make it challenging to predict how changes in input data or model parameters will affect the model's predictions. Consequently, biases or errors introduced during the training process may propagate and amplify, leading to biased and unfair outcomes.	Goodfellow, Bengio, and Courville (2016)
Feedback loops and reinforcement of biases	ML systems can inadvertently create feedback loops that perpetuate and amplify biases. For example, if a biased ML system is used to generate new data, this new data may also be biased, reinforcing the system's existing biases over time. These feedback loops can create a Butterfly Effect, where small initial biases lead to increasingly biased outcomes as the system iterates.	Ensign, Friedler, Neville, Scheidegger, and Venkatasubramanian (2018)
Compounding effects of multiple components	AI systems often comprise multiple components, each potentially introducing its biases. These biases can interact and compound in unpredictable ways, leading to a Butterfly Effect where the overall system exhibits greater bias than any single component.	Friedler, Scheidegger, and Venkatasubramanian (2016)
Local minima and distribution shifts	Saddle points in the loss landscape can stall optimization, causing algorithms to converge to suboptimal solutions that exacerbate fairness and bias issues. Distribution shifts, when test data differs from training data, can lead to poor model performance, unforeseen biases, and unfair outcomes. Both factors can trigger the Butterfly Effect on AI fairness and bias, amplifying minor discrepancies and yielding significant, unpredictable consequences.	Jordan (2023), Rezaei, Liu, Memarrast, and Ziebart (2021), Slowik and Bottou (2021)
Adversarial attacks	By exploiting small perturbations or vulnerabilities in AI models, adversarial attacks can intentionally introduce subtle changes to input data or manipulate the model's decision boundaries, causing the model to produce significantly different and biased outcomes. Small alterations introduced by adversarial attacks can lead to substantial and unpredictable consequences in AI fairness and bias, thereby manifesting the Butterfly Effect in these systems.	Nanda, Dooley, Singla, Feizi, and Dickerson (2021)

2. Examples of real-world emergence of the Butterfly Effect in AI systems

The Butterfly Effect can manifest in various ways in AI systems, leading to unintended consequences and exacerbating fairness and bias issues. Here, we discuss three real-world examples of the Butterfly Effect in AI systems.

2.1. Facial recognition technology

Facial recognition algorithms have become increasingly prevalent in various applications, from social media platforms to law enforcement. However, these algorithms can exhibit significant performance disparities across different demographic groups due to imbalanced training datasets. Small biases in the demographic representation of these datasets can lead to large differences in the algorithm's accuracy for different groups, resulting in biased outcomes that disproportionately affect underrepresented populations (Ezzeldin, Yan, He, Ferrara, & Avestimehr, 2023; Yan, Kao, & Ferrara, 2020).

In a study by Buolamwini and Gebru (2018), they found that commercial facial recognition systems had higher error rates for darker-skinned and female subjects compared to lighter-skinned and male subjects. These disparities in performance can be attributed to the underrepresentation of specific demographic groups in the training data, leading to a Butterfly Effect in the fairness and bias of facial recognition systems.

2.2. Healthcare algorithms

AI and ML models are increasingly being used to support decision-making in healthcare, such as identifying high-risk patients and guiding treatment decisions. However, biases in historical data and model assumptions can lead to biased predictions, disproportionately impacting minority populations.

Obermeyer, Powers, Vogeli, and Mullainathan (2019) found that a widely-used commercial algorithm for predicting healthcare needs exhibited significant racial bias. The algorithm assigned lower risk scores to Black patients than White patients with similar health conditions, leading to disparities in access to care management programs. The study revealed that the algorithm relied on healthcare costs as a proxy for health needs, which inadvertently introduced bias due to racial differences in healthcare utilization. The small initial bias in the model's assumptions led to a Butterfly Effect, resulting in large disparities in the allocation of healthcare resources.

2.3. Hiring algorithms

AI-based recruiting tools are increasingly being employed to streamline the hiring process and identify qualified candidates. However, these tools can perpetuate and amplify existing biases in the hiring process, leading to unfair outcomes and exacerbating societal inequalities (Raghavan, Barocas, Kleinberg, & Levy, 2020).

In 2018, it was reported that an AI recruiting tool showed gender bias, favoring male candidates over female candidates for technical roles (Dastin, 2018). The bias emerged due to the training data, which consisted primarily of resumes submitted to the company over a ten-year period and reflected a male-dominated applicant pool. Additionally, the algorithm might have penalized resumes containing words associated with women, such as "women's" in phrases like "women's chess club". The biases present in the training data and the use of gender-biased proxies led to a Butterfly Effect, resulting in a biased AI system that perpetuated gender inequality in hiring.

2.4. Large language models

Large language models, such as GPT-4 or Bard, could be significantly impacted by the Butterfly Effect. The training process for these models involves learning from vast amounts of text data, making them susceptible to minor changes in input data or algorithmic processes




Fig. 1. Root causes and associated manifestations of the Butterfly Effect in AI systems.

(Pernischová, 2019). A seemingly insignificant alteration in the training data, such as a slightly skewed representation of a particular demographic or viewpoint, can lead to substantial and unexpected biases in the model's output (Augenstein et al., 2023). Additionally, adversarial attacks, distributional shifts between training and test data could also lead to unforeseen outcomes. Hence, the Butterfly Effect could manifest in large language models, resulting in outputs that propagate and amplify pre-existing biases or inaccuracies — see Ferrara (2023b, 2023c), Weidinger et al. (2021).

3. Manifestations of the Butterfly Effect in AI systems

We explore the various manifestations of the Butterfly Effect on AI fairness and bias. These manifestations include small adjustments in input data, inherent biases within data or algorithms, feedback loops that amplify biases, and adversarial attacks exploiting vulnerabilities. By understanding these different ways in which the Butterfly Effect can impact AI systems, we can better identify potential sources of unfairness and bias and develop appropriate strategies to mitigate their effects, ensuring that AI systems produce fair and unbiased outcomes. Fig. 1 summarizes the causes and manifestations of the Butterfly Effect in AI systems presented in detail next.

3.1. Small adjustments in input data

Small adjustments in input data can significantly impact the fairness and bias of AI systems, as these systems rely on vast amounts of data to make decisions. These adjustments can manifest in various ways, including changes in data sampling, demographic makeup, and feature selection, among others. This sensitivity to input data can lead to the Butterfly Effect, where minor alterations in input data result in significant and unintended consequences in the fairness and bias of AI systems.

3.1.1. Data sampling

The manner in which data is collected and sampled can introduce biases in AI systems. If the data collection process is not carefully designed, it can lead to underrepresentation or overrepresentation of certain groups, which can in turn affect the fairness and bias of the AI system (Shankar et al., 2020). For example, non-representative sampling from a population can result in skewed datasets, causing AI systems to perform poorly for underrepresented groups.

3.1.2. Demographic makeup

AI systems may exhibit biases when demographic makeup within the training dataset is imbalanced (He & Garcia, 2009). For instance, if a particular demographic group is underrepresented in the training data, the AI system may not generalize well to that group, leading to biased and unfair outcomes (Buolamwini & Gebru, 2018). In facial recognition technology, the underrepresentation of certain demographic groups in training datasets can lead to disparities in the algorithm's accuracy for those groups. Ensuring that training datasets are representative of diverse populations is crucial to mitigate the Butterfly Effect and ensure the fairness of AI systems.

3.1.3. Feature selection and engineering

Feature selection and engineering play a crucial role in shaping the behavior of AI systems. The choice of features used as inputs can significantly affect the fairness and bias of AI models (Chouldechova, 2017). For example, using features that are proxies for protected attributes, such as race or gender, can introduce bias into AI systems, even if the protected attributes themselves are not explicitly used. Additionally, the omission of important features that capture relevant information about the population may result in biased models that do not adequately account for differences between groups.

In conclusion, small adjustments in input data, such as data sampling, demographic makeup, and feature selection, can have a profound impact on the fairness and bias of AI systems. To mitigate the Butterfly Effect and ensure that AI systems promote fairness and equity, it is essential to carefully curate and preprocess input data to minimize biases and accurately represent diverse populations.

3.2. Inherent biases within data or algorithms

We delve into the inherent biases within data or algorithms that can cause the Butterfly effect on AI fairness and bias. These biases can originate from various sources, such as biased data collection processes or unintentional biases embedded in algorithm design.

Inherent biases within data or algorithms can lead to the Butterfly Effect on AI fairness and bias. Data biases can arise from historical discrimination, measurement errors, or other systemic issues affecting the data-generating process, while algorithmic biases can emerge from model assumptions, optimization techniques, or other design choices. These biases can propagate and compound throughout the system, leading to significant and unintended consequences.

3.2.1. Biases in data

Data biases can emerge from various sources, such as historical discrimination or measurement errors. For instance, historical biases present in training data can lead AI systems to perpetuate or exacerbate existing inequalities. An example of this phenomenon is seen in the COMPAS recidivism risk assessment tool, which was found to have disparate impacts on different racial groups due to biases present in the training data (Angwin, Larson, Mattu, & Kirchner, 2016). Another source of data bias is measurement error, which occurs when variables in the dataset do not accurately represent the underlying constructs they are meant to capture. Measurement errors can introduce biases in AI systems, leading to unfair decision-making (Dressel & Farid, 2018).

3.2.2. Algorithmic biases

Algorithmic biases can arise from various aspects of the AI system, such as model assumptions, optimization techniques, or other design choices. Model assumptions, like the choice of a linear model or the assumption of independence between features, can introduce biases if they do not accurately reflect the underlying data-generating process (Berk, Heidari, Jabbari, Kearns, & Roth, 2018). Optimization techniques, such as regularization or the choice of a loss function, can also lead to biases if they prioritize certain objectives over others (Kearns, Neel, Roth, & Wu, 2018). Lastly, other design choices, like the selection

of hyperparameters or the choice of an ensemble method, can introduce biases into AI systems, which can propagate and compound over time (Grgić-Hlača, Zafar, Gummadi, & Weller, 2018).

3.3. Feedback loops that amplify biases

Feedback loops can amplify biases in AI systems, leading to the Butterfly Effect compromising fairness and bias. Feedback loops occur when the output of an AI system influences its future inputs, reinforcing and magnifying biases over time. This can result in a self-perpetuating cycle of unfair outcomes that disproportionately impact certain groups.

3.3.1. Reinforcing feedback loops

Reinforcing feedback loops can occur when an AI system's biased predictions lead to actions that further perpetuate the initial biases. For example, predictive policing algorithms that rely on historical crime data can create a feedback loop by directing law enforcement resources to areas with higher reported crime rates (Lum & Isaac, 2016). If certain groups or neighborhoods are disproportionately targeted due to historical biases in the data, the increased police presence can lead to more arrests and crime reports, which in turn reinforce the initial biases in the AI system.

3.3.2. Feedback loops in recommendation systems

Recommendation systems are another example where feedback loops can amplify biases. These systems often rely on user data to provide personalized recommendations, which can create a filter bubble that reinforces users' pre-existing preferences and biases (Nguyen, Hui, Harper, Terveen, & Konstan, 2014). In turn, this can lead to biased content exposure and a lack of diversity in the information that users are exposed to, perpetuating existing social biases and contributing to polarization (Pariser, 2011).

3.3.3. Algorithmic confounding

Algorithmic confounding is a phenomenon where a feedback loop between an AI system's predictions and the ground truth it aims to predict leads to biased and unfair outcomes. This can occur when the AI system's biased predictions influence the data used to evaluate its performance, making it difficult to disentangle the true effect of the AI system from the biases present in the data. In such cases, biased algorithms may appear to perform well due to the confounding effect, reinforcing the initial biases and leading to a self-perpetuating cycle of unfair outcomes.

To address feedback loops and mitigate their amplifying effect on biases, it is essential to carefully consider the potential consequences of AI systems' predictions on their future inputs and design mechanisms for monitoring and correcting biases as they emerge over time.

3.4. Adversarial attacks exploiting vulnerabilities

Adversarial attacks can exploit vulnerabilities in AI systems, leading to the Butterfly effect that compromises fairness and bias. These attacks involve the intentional manipulation of input data or model parameters to cause an AI system to produce biased, unfair, or otherwise undesirable outcomes. By exploiting the sensitivity of AI systems to small changes, adversarial attacks can have significant and unpredictable consequences on fairness and bias.

3.4.1. Adversarial examples

Adversarial examples are carefully crafted input data designed to cause an AI system to produce incorrect or biased predictions (Szegedy et al., 2014). These examples can be generated by adding small, imperceptible perturbations to the input data, which can result in vastly different predictions due to the Butterfly Effect. Adversarial examples can be particularly problematic for fairness and bias, as they can be used to target specific demographic groups or individuals, leading to discriminatory outcomes (Szegedy et al., 2014).

3.4.2. Model inversion and membership inference attacks

Model inversion and membership inference attacks can exploit the vulnerabilities of AI systems to reveal sensitive information about the training data or individuals within the dataset (Fredrikson, Jha, & Ristenpart, 2015; Shokri, Stronati, Song, & Shmatikov, 2017). These attacks can lead to the Butterfly Effect compromising fairness and bias by revealing disparities in the demographic makeup of the training data or exposing biases in the AI system's predictions. Furthermore, the knowledge gained from these attacks can be used to craft more sophisticated adversarial examples, further exacerbating the impact of the Butterfly Effect on fairness and bias.

3.4.3. Poisoning attacks

Poisoning attacks involve the manipulation of the training data or model parameters to introduce or amplify biases in an AI system (Solans, Biggio, & Castillo, 2020). By injecting carefully crafted examples into the training data or modifying the model parameters, adversaries can exploit the Butterfly Effect to cause an AI system to produce biased, unfair, or otherwise undesirable outcomes. Poisoning attacks can be particularly challenging to detect and mitigate, as the perturbations introduced by the attacker can be subtle and difficult to distinguish from natural variations in the data.

To defend against adversarial attacks and mitigate their impact on fairness and bias, it is essential to develop robust AI systems that can withstand small perturbations in the input data or model parameters (Madry, Makelov, Schmidt, Tsipras, & Vladu, 2017). Additionally, regular monitoring and evaluation of AI systems for fairness and bias, as well as the implementation of privacy-preserving techniques, can help prevent adversaries from exploiting the Butterfly Effect to compromise the fairness and integrity of AI systems.

4. Strategies to mitigate the butterfly effect on AI fairness and bias

Next, we discuss various strategies to mitigate the Butterfly Effect on AI fairness and bias. These strategies encompass diverse aspects of AI system development, ranging from data collection and preprocessing to algorithmic fairness, evaluation and monitoring, and adversarial robustness. By employing these strategies, researchers and practitioners can work towards addressing the potential unintended consequences arising from small changes in input data or algorithmic design and ensure that AI systems are more transparent, accountable, and fair. Fig. 2 summarizes the mitigation strategies presented in detail next.

4.1. Data collection and preprocessing

Creating balanced and representative datasets is crucial for mitigating the Butterfly Effect on AI fairness and bias. Several techniques can be employed to ensure that datasets are balanced and accurately represent the population of interest.

1. **Oversampling minority classes:** Oversampling involves creating copies of instances from minority classes to balance the class distribution. One well-known technique is the Synthetic Minority Over-sampling Technique (SMOTE), which generates synthetic instances of minority classes by interpolating between existing instances (Chawla, Bowyer, Hall, & Kegelmeyer, 2002). SMOTE can help alleviate the problem of overfitting associated with simple oversampling and lead to better performance in terms of fairness and generalization.
2. **Undersampling majority classes:** Undersampling involves removing instances from majority classes to balance the class distribution. One effective undersampling technique is Tomek Links, which removes majority class instances that are close to the decision boundary (Kubat & Matwin, 1997). By removing these instances, the decision boundary becomes less sensitive to small changes in the data, mitigating the Butterfly Effect.




Fig. 2. Strategies to mitigate the Butterfly Effect on AI fairness and bias.

3. **Synthetic data generation:** Synthetic data generation can be used to create new, artificial instances for underrepresented classes, ensuring that the dataset is representative of the population of interest. Techniques such as Variational Autoencoders (VAEs) and Generative Adversarial Networks (GANs) can generate high-quality synthetic data that closely resembles the real-world distribution of the data, reducing the impact of small changes in the data on AI fairness and bias (Frid-Adar et al., 2018).
4. **Stratified sampling:** Stratified sampling is a method of sampling that involves dividing the population into homogeneous groups, or strata, and sampling instances from each stratum in proportion to the size of the stratum. This technique can help ensure that the dataset is representative of the population of interest and reduce the sensitivity of AI systems to small changes in the data.

4.2. Algorithmic fairness

Algorithmic fairness is a critical aspect of machine learning that focuses on ensuring equitable treatment across different groups, mitigating biases, and reducing the sensitivity of models to minor changes in input data.

1. **Fairness-aware machine learning:** Fairness-aware machine learning aims to incorporate fairness constraints during the training process to minimize disparate treatment and disparate impact on different groups. Zafar, Valera, Gomez-Rodriguez, and Gummadi (2017) propose a convex optimization formulation for learning a classifier that satisfies various notions of fairness while maintaining high accuracy. The method minimizes the difference in correlations between the classifier's predictions and the sensitive attribute (e.g., race, gender) across different groups. By incorporating fairness constraints during the training process, this approach helps to mitigate the Butterfly Effect by reducing the sensitivity of the model to small changes in the input data.
2. **Post-processing methods for achieving fairness:** Post-processing techniques focus on adjusting the output of a trained model to ensure fairness. Hardt, Price, and Srebro (2016) propose a method that learns a transformation of the classifier's predictions to satisfy the equalized odds criterion (equal true positive and false positive rates across groups). The method requires no retraining of the original classifier and guarantees the best achievable trade-off between accuracy and fairness. By adjusting the model's output, post-processing methods can help to mitigate the Butterfly Effect by compensating for biases that may have been introduced during the training process.

3. **Fairness through awareness:** Dwork, Hardt, Pitassi, Reingold, and Zemel (2012) introduce the concept of fairness through awareness, which requires that any two individuals who are similar with respect to a specific task should be treated similarly by the algorithm. They propose a Lipschitz condition on the classifier's behavior, ensuring that the model is less sensitive to small changes in the input data. Fairness through awareness can help to mitigate the Butterfly Effect by constraining the model's behavior to be consistent and less influenced by small perturbations in the data.
4. **Regularization for fairness:** Regularization techniques can be used to encourage fairness in AI models. Zhao et al. (2019) propose a method that incorporates a fairness-aware regularization term in the objective function during training. This term penalizes the difference between the predicted outcomes for different groups, encouraging the model to learn more equitable representations of the data. Regularization for fairness can help to mitigate the Butterfly Effect by discouraging the model from relying on small differences in the input data that may lead to unfair outcomes.

4.3. Evaluation and monitoring

We emphasize the importance of assessing, understanding, and continuously overseeing the fairness and behavior of AI systems to detect and address potential biases and unintended consequences due to the Butterfly Effect.

1. **Fairness-aware performance metrics:** Evaluating the fairness of AI models requires metrics that capture the disparate impact on different groups. Verma and Rubin (2018) discuss a set of fairness-aware performance metrics, such as demographic parity, equalized odds, and equal opportunity. These metrics provide a quantitative measure of the difference in outcomes for different groups, helping to identify potential bias and unfairness in AI models. By using fairness-aware performance metrics, practitioners can monitor the impact of small changes in the data or model and identify instances where the Butterfly Effect leads to unintended consequences.
2. **Auditing tools for AI systems:** Auditing tools can help to systematically analyze the behavior of AI systems to identify potential biases, fairness violations, and other issues. Grgić-Hlača et al. (2018) present a method for auditing black-box models that involves perturbing the input data to explore the model's sensitivity to different features, specifically focusing on protected attributes. By systematically analyzing the model's behavior, auditing tools can help to identify potential Butterfly Effects and inform subsequent mitigation strategies.

3. **Model interpretability for fairness:** Model interpretability techniques can provide insights into the decision-making process of AI models, allowing for better scrutiny of potential bias and unfairness. Ribeiro, Singh, and Guestrin (2016) propose Local Interpretable Model-agnostic Explanations (LIME), a method for explaining the predictions of any classifier by approximating it locally with an interpretable model. By understanding the underlying factors that influence a model's decisions, practitioners can detect potential Butterfly Effects and address them through appropriate interventions.
4. **Continual monitoring and feedback:** Continual monitoring and feedback involve tracking the performance and fairness of AI systems over time, as well as collecting feedback from users to identify potential issues. Mitchell et al. (2019) discuss the importance of actively soliciting user feedback to identify biases, unfairness, and other issues that may not be apparent from traditional evaluation metrics. Continual monitoring and feedback can help to uncover Butterfly Effects that emerge over time or as a result of changing data distributions and user contexts.

4.4. Adversarial robustness

We delve into the challenges and solutions associated with ensuring AI models remain fair and unbiased when faced with adversarial attacks.

1. **Robustness bias:** Nanda et al. (2021) suggested that relying solely on traditional notions of fairness based on a model's outputs may not be enough when models are susceptible to adversarial attacks. To measure robustness bias, they proposed two methods and performed an empirical investigation on state-of-the-art deep neural networks using commonly employed real-world fairness datasets. Their findings demonstrated that subgroups, which are categorized by sensitive attributes like race and gender, are less robust and therefore more susceptible to adversarial attacks. It is important to consider robustness bias when evaluating real-world systems that rely on deep neural networks to make decisions (Nanda et al., 2021).
2. **Techniques for improving AI model robustness:** To mitigate the Butterfly Effect on AI fairness and bias caused by adversarial attacks, it is essential to develop models that are robust against such perturbations. Madry et al. (2017) propose a training framework based on adversarial training, which aims to minimize the worst-case loss over a set of allowed perturbations. By training the model to be robust against adversarial examples, this approach helps to improve the model's robustness against small changes in the input data, reducing the risk of unintended consequences due to the Butterfly Effect.
3. **Defense strategies against adversarial attacks:** Defending against adversarial attacks is crucial for ensuring the robustness of AI models and mitigating the Butterfly Effect. Tramèr, Kurakin, et al. (2017), Tramèr, Papernot, Goodfellow, Boneh, and McDaniel (2017) propose a defense strategy called ensemble adversarial training, which augments the training data with adversarial examples generated by an ensemble of models. This approach helps to increase the diversity of the adversarial examples used during training, making the model more robust against a broader range of attacks. By defending against adversarial attacks, this approach helps to reduce the impact of the Butterfly Effect on AI fairness and bias.
4. **Certified robustness:** Certified robustness aims to provide guarantees on the model's behavior under adversarial perturbations, ensuring that small changes in the input data do not lead to significant changes in the model's output. Cohen, Rosenfeld, and Kolter (2019) introduce randomized smoothing, a technique that provides provable robustness guarantees for classifiers against ℓ_2 -norm bounded adversarial perturbations. By providing certified robustness, this approach helps to mitigate the Butterfly Effect by ensuring that the model's behavior is stable under small perturbations in the input data.
5. **Adversarial detection:** Detecting adversarial examples is an essential step in defending against adversarial attacks and mitigating the Butterfly Effect on AI fairness and bias. Metzen, Genewein, Fischer, and Bischoff (2017) propose a method for detecting adversarial examples by training a separate neural network to distinguish between clean and adversarial inputs. By detecting adversarial examples, this approach helps to prevent unintended consequences due to the Butterfly Effect, ensuring that the model's behavior remains fair and unbiased.

5. Conclusions

In this section, we synthesize the insights and observations presented in our exploration of the Butterfly Effect within AI fairness and bias. Beginning with a detailed discussion on the nuanced impacts of small changes in AI systems (Section 5.1), we delve into the broader societal implications these changes harbor, emphasizing the critical need for ethically aligned AI development (Section 5.2). The section culminates in outlining future directions, proposing enhanced strategies for mitigating bias, advocating for interdisciplinary collaboration, and emphasizing the necessity of adaptable, user-centric AI systems (Section 5.3). We hope to not only encapsulate our findings but also chart a course for future research and policy-making in the realm of responsible AI.

5.1. Discussion

In exploring the Butterfly Effect in AI fairness and bias, this paper delved into the intricate dynamics of small changes within AI systems and their potential disproportionate impacts. It underscored the interconnectedness and complexity of AI systems, highlighting the necessity for AI practitioners to rigorously assess even minor modifications. This discussion aims to deepen the understanding of how subtle nuances in data and algorithms can lead to significant ethical and functional shifts in AI outputs.

The ripple effects of small changes. We elucidated the intricate and often unforeseen consequences of minor alterations in AI systems, characterized by the Butterfly Effect in AI fairness and bias. It is pivotal to recognize how these small changes, whether in data input, algorithmic adjustments, or even subtle shifts in data distribution, can catalyze far-reaching and unpredictable impacts on the fairness and bias of AI systems. The significance of this understanding lies in its ability to guide AI practitioners in scrutinizing every minor change in their systems, acknowledging that even negligible alterations can yield disproportionate outcomes, especially for underrepresented groups.

Interconnectedness and complexity of AI systems. The multifaceted nature of AI systems demands a comprehensive approach to examine the Butterfly Effect. The interplay of various components within these systems, from data sampling and algorithmic biases to feedback loops and adversarial attacks, creates a complex web of interactions. This complexity underscores the need for holistic strategies that address the Butterfly Effect, ensuring that AI systems are not only technically proficient but also ethically sound and socially responsible.

5.2. Implications

We examined the far-reaching societal implications of the Butterfly Effect in AI. We reflected on the critical need for incorporating fairness and bias considerations into the very fabric of AI development. The implications extend beyond technical realms, advocating for robust policies and governance frameworks that can effectively oversee and regulate AI systems, ensuring they adhere to ethical standards and societal values.

Societal impact. The Butterfly Effect in AI has profound societal implications. As AI systems become more pervasive in critical domains like healthcare, law enforcement, and employment, the potential for small biases to cause significant harm escalates. This scenario necessitates a vigilant approach to AI development, where fairness and bias are not afterthoughts but integral elements of the design and implementation process.

Policy and governance. This phenomenon also highlights the need for robust policies and governance structures that prioritize fairness in AI. Regulatory frameworks should mandate thorough evaluation of AI systems for potential Butterfly Effects, ensuring that AI developers and users remain accountable for the societal impact of these technologies.

5.3. Future directions

Focusing on the path ahead, this paper outlines the pivotal areas for future research and development in addressing the Butterfly Effect in AI. We call for enhanced mitigation strategies, emphasizing the value of interdisciplinary collaboration, and stressing the importance of continuous learning and adaptation in AI systems. This future outlook envisions AI development that is not only technically advanced but also ethically sound, user-centric, and socially responsible.

Enhanced mitigation strategies. Future research should focus on developing advanced methodologies to detect, quantify, and mitigate the Butterfly Effect in AI. This includes the creation of more sophisticated tools for monitoring AI systems in real-time, ensuring that they remain fair and unbiased throughout their lifecycle.

Collaboration and interdisciplinary approaches. Addressing the Butterfly Effect in AI requires a collaborative effort, drawing expertise from fields like computer science, social science, ethics, and law (Mittelstadt, Allo, Taddeo, Wachter, & Floridi, 2016). Interdisciplinary research can provide more nuanced insights into how AI systems interact with societal structures and human behavior, leading to more effective mitigation strategies.

Continuous learning and adaptation. AI systems should be designed with the capability to learn and adapt continuously, ensuring they remain resilient to the dynamic nature of data and societal contexts (Park et al., 2023). This adaptability is key to preventing the amplification of biases and ensuring long-term fairness in AI applications.

Focus on user-centric design. Finally, future AI systems should be developed with a strong focus on user-centric design, ensuring that they cater to the diverse needs and contexts of their users. This approach can help in creating AI systems that are not only fair and unbiased but also truly beneficial and inclusive for all segments of society.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

No data was used for the research described in the article.

References

Angwin, J., Larson, J., Mattu, S., & Kirchner, L. (2016). Machine bias: There's software used across the country to predict future criminals. And it's biased against blacks. *Augenstein, I., Baldwin, T., Cha, M., Chakraborty, T., Ciampaglia, G. L., Corney, D., DiResta, R., Ferrara, E., Hale, S., Halevy, A., Hovy, E., et al. (2023). Factuality challenges in the era of large language models. arXiv:2310.05189.*

Barocas, S., & Selbst, A. D. (2016). Big data's disparate impact. *California Law Review*, 104(671).

Berk, R., Heidari, H., Jabbari, S., Kearns, M., & Roth, A. (2018). Fairness in criminal justice risk assessments: The state of the art. *Sociological Methods & Research*, 50(1).

Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. In *Proceedings of the 1st conference on fairness, accountability and transparency* (pp. 77–91).

Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16, 321–357.

Chouldechova, A. (2017). Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5(2), 153–163.

Cohen, J. M., Rosenfeld, E., & Kolter, J. Z. (2019). Certified adversarial robustness via randomized smoothing. In *Proceedings of the 36th international conference on machine learning* (pp. 1310–1320).

Dastin, J. (2018). Amazon scraps secret AI recruiting tool that showed bias against women. *Reuters*.

Dressel, J., & Farid, H. (2018). The accuracy, fairness, and limits of predicting recidivism. *Science Advances*, 4(1), eaao5580.

Dworki, C., Hardt, M., Pitassi, T., Reingold, O., & Zemel, R. (2012). Fairness through awareness. In *Proceedings of the 3rd innovations in theoretical computer science conference* (pp. 214–226).

Ensign, D., Friedler, S. A., Neville, S., Scheidegger, C., & Venkatasubramanian, S. (2018). Runaway feedback loops in predictive policing. In *Proceedings of machine learning research*, vol. 81 (pp. 1–12).

Ezzeldin, Y. H., Yan, S., He, C., Ferrara, E., & Avestimehr, A. S. (2023). Fairfed: Enabling group fairness in federated learning. In *Proceedings of the AAAI conference on artificial intelligence*, vol. 37, no. 6 (pp. 7494–7502).

Ferrara, E. (2023a). Fairness and bias in artificial intelligence: A brief survey of sources, impacts, and mitigation strategies. *arXiv:2304.07683*.

Ferrara, E. (2023b). GenAI against humanity: Nefarious applications of generative artificial intelligence and large language models. *arXiv:2310.00737*.

Ferrara, E. (2023c). Should ChatGPT be biased? Challenges and risks of bias in large language models. *First Monday*, 28(11).

Fredrikson, Matt, Jha, Somesh, & Ristenpart, Thomas (2015). Model inversion attacks that exploit confidence information and basic countermeasures. In *Proceedings of the 22nd ACM SIGSAC conference on computer and communications security* (pp. 1322–1333).

Frid-Adar, M., Diamant, I., Klang, E., Amitai, M., Goldberger, J., & Greenspan, H. (2018). GAN-based synthetic medical image augmentation for increased CNN performance in liver lesion classification. *Neurocomputing*, 321, 321–331.

Friedler, S. A., Scheidegger, C., & Venkatasubramanian, S. (2016). On the (im) possibility of fairness. *arXiv:1609.07236*.

Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press.

Grgić-Hlača, N., Zafar, M. B., Gummadi, K. P., & Weller, A. (2018). Beyond distributive fairness in algorithmic decision making: Feature selection for procedurally fair learning. In *Proceedings of the thirty-second AAAI conference on artificial intelligence*. In *Advances in neural information processing systems* (pp. 3315–3323).

Hardt, M., Price, E., & Srebro, N. (2016). Equality of opportunity in supervised learning. In *Advances in neural information processing systems* (pp. 3315–3323).

He, H., & Garcia, E. A. (2009). Learning from imbalanced data. *IEEE Transactions on Knowledge and Data Engineering*, 21(9), 1263–1284.

Jordan, K. (2023). Calibrated chaos: Variance between runs of neural network training is harmless and inevitable. *arXiv:2304.01910*.

Kearns, M., Neel, S., Roth, A., & Wu, Z. S. (2018). Preventing fairness gerrymandering: Auditing and learning for subgroup fairness. In *Proceedings of the 35th international conference on machine learning* (pp. 2564–2572).

Kubat, M., & Matwin, S. (1997). Addressing the curse of imbalanced training sets: one-sided selection. In *Proceedings of the fourteenth international conference on machine learning* (pp. 179–186).

Lorenz, E. N. (1963). Deterministic nonperiodic flow. *Journal of Atmospheric Sciences*, 20(2), 130–141.

Lum, K., & Isaac, W. (2016). To predict and serve? *Significance*, 13(5), 14–19.

Madry, A., Makelov, A., Schmidt, L., Tsipras, D., & Vladu, A. (2017). Towards deep learning models resistant to adversarial attacks. In *Proceedings of the 6th international conference on learning representations*.

Metzen, J. H., Genewein, T., Fischer, V., & Bischoff, B. (2017). On detecting adversarial perturbations. *arXiv:1702.04267*.

Mitchell, M., Wu, S., Zaldívar, A., Barnes, P., Vasserman, L., Hutchinson, B., ... Gebru, T. (2019). Model cards for model reporting. In *Proceedings of the conference on fairness, accountability, and transparency* (pp. 220–229).

Mittelstadt, B. D., Allo, P., Taddeo, M., Wachter, S., & Floridi, L. (2016). The ethics of algorithms: Mapping the debate. *Big Data & Society*, 3(2), Article 2053951716679679.

Nanda, V., Dooley, S., Singla, S., Feizi, S., & Dickerson, J. P. (2021). Fairness through robustness: Investigating robustness disparity in deep learning. In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency* (pp. 466–477).

Nguyen, T. T., Hui, P.-M., Harper, F. M., Terveen, L., & Konstan, J. A. (2014). Exploring the filter bubble: The effect of using recommender systems on content diversity. In *Proceedings of the 23rd international conference on world wide web* (pp. 677–686).

Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464), 447–453.

- Pariser, E. (2011). *The filter bubble: What the internet is hiding from you*. Penguin UK.
- Park, J. S., O'Brien, J. C., Cai, C. J., Morris, M. R., Liang, P., & Bernstein, M. S. (2023). Generative agents: Interactive simulacra of human behavior. arXiv:2304.03442.
- Pernischová, R. (2019). *The butterfly effect in knowledge graphs: Predicting the impact of changes in the evolving web of data*. ISWC.
- Raghavan, M., Barocas, S., Kleinberg, J., & Levy, K. (2020). Mitigating bias in algorithmic hiring: Evaluating claims and practices. In *Proceedings of the 2020 conference on fairness, accountability, and transparency* (pp. 469–481).
- Rezaei, A., Liu, A., Memarrast, O., & Ziebart, B. D. (2021). Robust fairness under covariate shift. In *Proceedings of the AAAI conference on artificial intelligence*, vol. 35, no. 11 (pp. 9419–9427).
- Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). Why should I trust you?: Explaining the predictions of any classifier. In *Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining* (pp. 1135–1144).
- Shankar, S., Halpern, Y., Breck, E., Atwood, J., Wilson, J., & Sculley, D. (2020). No classification without representation: Assessing geodiversity issues in open data sets for the developing world. In *Proceedings of the 2020 conference on fairness, accountability, and transparency* (pp. 445–454).
- Shokri, Reza, Stronati, Marco, Song, Congzheng, & Shmatikov, Vitaly (2017). Membership inference attacks against machine learning models. In *2017 IEEE symposium on security and privacy (SP)* (pp. 3–18). IEEE.
- Stowik, A., & Bottou, L. (2021). Algorithmic bias and data bias: Understanding the relation between distributionally robust optimization and data curation. arXiv: 2106.09467.
- Solans, David, Biggio, Battista, & Castillo, Carlos (2020). Poisoning attacks on algorithmic fairness. In *Joint European conference on machine learning and knowledge discovery in databases* (pp. 162–177). Springer.
- Strogatz, S. H. (2018). *Nonlinear dynamics and chaos with student solutions manual: With applications to physics, biology, chemistry, and engineering*. CRC Press.
- Szegedy, Christian, Zaremba, Wojciech, Sutskever, Ilya, Bruna, Joan, Erhan, Dumitru, Goodfellow, Ian, & Fergus, Rob (2014). Intriguing properties of neural networks. In *2nd international conference on learning representations*.
- Tramèr, F., Kurakin, A., Papernot, N., Goodfellow, I., Boneh, D., & McDaniel, P. (2017). Ensemble adversarial training: Attacks and defenses. In *Proceedings of the 6th international conference on learning representations*.
- Tramèr, F., Papernot, N., Goodfellow, I., Boneh, D., & McDaniel, P. (2017). The space of transferable adversarial examples. arXiv:1704.03453.
- Verma, S., & Rubin, J. (2018). Fairness definitions explained. In *Proceedings of the 2018 ACM conference on international computing education research* (pp. 121–129).
- Weidinger, L., Mellor, J., Rauh, M., Griffin, C., Uesato, J., Huang, P. S., Gabriel, I. (2021). Ethical and social risks of harm from language models. arXiv:2112.04359.
- Yan, S., Kao, H.-t., & Ferrara, E. (2020). Fair class balancing: Enhancing model fairness without observing sensitive attributes. In *Proceedings of the 29th ACM international conference on information & knowledge management* (pp. 1715–1724).
- Zafar, M. B., Valera, I., Gomez-Rodriguez, M., & Gummadi, K. P. (2017). Fairness beyond disparate treatment & disparate impact: Learning classification without disparate mistreatment. In *Proceedings of the 26th international conference on world wide web* (pp. 1171–1180).
- Zhao, J., Wang, T., Yatskar, M., Cotterell, R., Ordonez, V., & Chang, K. W. (2019). Gender bias in contextualized word embeddings. In *Proceedings of NAACL-HLT* (pp. 629–634).