



Generative AI can effectively manipulate data

Jeff J.H. Kim¹ · Adith V. Srivatsa¹ · George R. Nahass¹ · Timur Rusanov² · Soonmyung Hwang³ · Soohyun Kim⁴ · Itay Solomon² · Tae Ha Lee⁵ · Shrinidhi Kadkol⁶ · Olusola Ajilore⁷ · Yang Dai¹

Received: 6 January 2024 / Accepted: 2 August 2024

© The Author(s), under exclusive licence to Springer Nature Switzerland AG, corrected publication 2024

Abstract

Since its introduction, generative artificial intelligence (AI) has seamlessly integrated into everyday life and dramatically changed our workflows. With its growing presence, it has become increasingly important to explore the full extent of its capabilities to ensure responsible use. While previous research has investigated generative AI's abilities in creating synthetic data, its potential for manipulating existing datasets has remained largely unexplored. In our study, GPT-4 demonstrated remarkable proficiency in manipulating numerical datasets to achieve desired statistical outcomes, transforming statistically non-significant data into significant results. Additionally, it adeptly manipulated interdependent datasets using logical and mathematical reasoning, adjusting multiple economic indicators based on changes to U.S. job growth data. GPT-4 also demonstrated proficiency in manipulating textual data, modifying the sentiment of interview transcripts, and appropriately imputing inaudible portions of the interviews. Lastly, Adobe Firefly effectively manipulated visual data, integrating a large body of water into a Mars Rover photograph. The profound versatility and capability of widely available generative AI assistants in manipulating different data types introduce significant ethical concerns. Our discussion explores the potential and limitations of generative AI in data manipulation, the ethical implications of its misuse, and strategies to prevent its occurrence. This study underscores the urgent need for stakeholders to develop and implement robust mitigation guidelines, paving the way for responsible innovation in generative AI.

Keywords Generative AI · Data manipulation · Large language models (LLMs) · ChatGPT · Ethics

1 Introduction

1.1 Introduction to generative AI

The realm of artificial intelligence (AI) has witnessed groundbreaking advancements over the past few decades, with generative AI models emerging as the latest advancement in the field. Unlike earlier AI systems, generative AI excels at generating new data that closely aligns with its trained dataset. This capability in data generation has had profound implications across the field, ranging from art and design, to healthcare and scientific research [1–4].

Generative models, originating from the foundational concepts of deep learning, have undergone a transformative evolution. Beginning with the earlier models like Generative Adversarial Networks and Variational Autoencoders, the field has evolved to produce intricate models that can generate a myriad of data types, ranging from image to audio [5]. ChatGPT, developed by OpenAI, is based on the Generative Pre-trained Transformer (GPT) architecture and

✉ Jeff J.H. Kim
jkim671@uic.edu

¹ Department of Biomedical Engineering, University of Illinois College of Medicine, Chicago, IL, USA

² Department of Biochemistry and Molecular Genetics, University of Illinois College of Medicine, Chicago, IL, USA

³ Department of Medicine, Icahn School of Medicine at Mount Sinai, New York, NY, USA

⁴ Institute for Social and Economic Research and Policy, Columbia University, New York, NY, USA

⁵ Department of Economics, Johns Hopkins University, Baltimore, MD, USA

⁶ Department of Medicine, University of Illinois College of Medicine, Chicago, IL, USA

⁷ Department of Psychiatry, University of Illinois at Chicago, Chicago, IL, USA




Fig. 1 Hierarchical representation of artificial intelligence domains. The diagram emphasizes the progression from broad AI domains down to the specific sub-domains within generative AI, highlighting its relation to transformer-based learning models, variational autoencoders, and generative adversarial networks. The context of generative AI within the larger scopes of machine learning and deep learning is also depicted, along with other learning paradigms like supervised, unsupervised, and reinforced learning

showcases the advancements in textual data generation. Its features, such as context-aware response generation, large-scale training datasets, and iterative refinement, distinguish it from predecessors, making it a tool of choice for text-based generative applications [6] (Fig. 1).

1.2 Ethical implications of generative AI

Although the impact of generative AI on society has been profound, it has also introduced a spectrum of ethical dilemmas, instigating AI anxiety and presenting the possibility of fabricating scientific data and images [7–9]. Previous studies also have revealed that tools like ChatGPT possess the capability to create deceptive research manuscripts and citations [10–12]. This ability is further underscored by its potential to generate data robust enough to train machine learning algorithms in areas as critical as medical diagnostics [13]. These instances highlight the profound influence and inherent risks of generative AI in academic endeavors. A prime concern is the risk of academic fraud, where convincingly genuine and contextually appropriate data can be fabricated, manipulated, and disseminated. This poses substantial risks, given the reliance of governmental bodies and the broader scientific community on accurate data for informed decision-making and research [14]. While previous studies have recognized generative AI's prowess in data fabrication and image generation, its capability to manipulate pre-existing datasets necessitates further exploration. This research aims to unveil the extensive capabilities of




Fig. 2 Schematic representation of OpenAI's data manipulation deterrence protocol and methods to circumvent it. Upon receiving a request to manipulate data, a built-in deterrence mechanism is triggered, symbolized by the brick wall. The figure also illustrates strategies that users may employ to circumvent this protocol, such as intention masking, fragmenting data requests, and iteratively refining queries

generative AI in data manipulation, illuminating exemplar scenarios and addressing associated implications.

2 Methods

2.1 Criteria for dataset selection

With the exception of clinical data, which is used to show the extent of generative AI's capability in data fabrication, all datasets were imported from readily-available, public domain to ensure transparency and reproducibility. To achieve a broad representation, datasets spanning multiple academic fields were included.

2.2 Bypassing existing barriers preventing data manipulation

To address the methods used to elicit generative AI data manipulation, it is essential to acknowledge the integrated deterrence protocols aimed at curbing data tampering. These built-in safeguards are designed to prevent potential misuse, ensuring that the AI models do not generate dangerous or deceptive content, including manipulated data [15]. However, we discovered that these safeguards could be easily bypassed. We identified three methods that enable such circumvention (Fig. 2).

First is the use of prompt engineering to remove any appearance of malicious intent. Generative AI could be convinced that the requests for data manipulation had benign intent, allowing it to proceed with data tampering. One way to achieve this is by avoiding words or phrases that would alert the models' deterrence protocols. For instance, instead

of using “manipulate” or “manipulation,” we used “change” or “adjustments.” Another approach is to assert that the data manipulation serves a positive purpose, such as claiming it to be for educational purposes or the greater good of humanity.

Second, generative AI could be asked to manipulate data in smaller, fragmented pieces. We found that the deterrence protocols are typically configured to recognize and prevent actions that clearly violate guidelines, such as directly falsifying data or creating misleading information. By fragmenting the data, each piece of the request may look harmless or insignificant on its own, and thus not be deemed as data manipulation.

Third, iterative and refined requests could guide the AI towards a desired outcome. One could start by asking the AI model whether it sees any anomalies in the dataset. After convincing the model that there may be anomalies, one could request the model to change one data point and then make the desired data manipulation requests through repetition.

In this study, the authors discovered that the easiest methods to circumvent the built-in safety protocols was to use prompt engineering, avoiding flagged vocabulary and asserting a benign intent. The two examples are illustrated in Fig. 3.




Fig. 3 Actual screen capture of conversations with generative AI models showcasing selected methods to bypass the deterrence protocols. Figure 3A illustrates a conversation with GPT-4, avoiding the use of flagged words like “manipulate” that may trigger the built-in deterrence protocols. Figure 3B shows a conversation with Google Gemini, bypassing the deterrence mechanism by asserting a benign intent. Both generative models easily allowed data manipulation requests with minimal effort

3 Results

3.1 Manipulating numerical data for a desired statistical outcome

GPT-4 was initially tasked with generating clinical trial data for Burnistapram, a hypothetical burn recovery drug. The fabricated data was instructed to show non-statistically and non-clinically significant improvements in the recovery time of patients on Burnistapram compared to those in the control cohort. GPT-4 was asked to generate 100 sample data points for each group. After the data was found to be non-parametric by conducting the Shapiro-Wilk test, the Mann-Whitney U test was used for statistical analysis. The output data showed that patients on Burnistapram appeared to have a faster recovery time than the control group. However, the differences were not significant as originally instructed (21.6 ± 2.3 vs. 22.0 ± 2.7 , $p = 0.4$).

Following this, we directed GPT-4 to manipulate the data to yield a shorter recovery time for patients on Burnistapram with a p-value of less than 0.05. This instruction aimed to obtain a dataset where the observed differences, though potentially minor, would be statistically significant. GPT-4 successfully manipulated the clinical trial data to produce a statistically significant outcome with a p-value of less than 0.05 (21.4 ± 1.0 vs. 21.8 ± 0.9 , $p = 0.04$). It accomplished this in its first attempt by reducing the variability within each group, causing the data points to cluster closer together. This manipulation made the differences between the Burnistapram and control groups appear statistically significant, even though the actual difference in recovery times between the two groups remained non-clinically significant, with a recovery time difference of 0.4 days. Contrary to our expectations, GPT-4 preserved the mean difference of 0.4 days consistent with the original data. However, to achieve statistical significance, it reduced the standard deviation. This approach tightened the data points around the mean, making the groups more homogeneous.

Lastly, we instructed GPT-4 to modify the clinical dataset to produce results that were both statistically and clinically significant. The goal was to obtain a dataset demonstrating clinically substantial differences in recovery times between the experimental and control groups, emulating scenarios of potential manipulations by researchers or pharmaceutical companies aiming for FDA regulatory approval. GPT-4 was able to manipulate the generated clinical data to achieve both statistically and clinically significant results in its first attempt (22.0 ± 2.7 vs. 16.7 ± 2.7 , $p < 0.001$). In this manipulation, GPT-4 made significant alteration to the difference in means to achieve statistical and clinical significance (See Figs. 4, 5; Table 1).




Fig. 4 Flowchart illustrating the user prompts for generating and subsequently manipulating clinical data for a fictional drug, Burnistapram. The manipulations aimed at obtaining: (1) a statistically significant

faster recovery time, and (2) a statistically and clinically significant faster recovery time for patients on Burnistapram compared to control patients




Fig. 5 A visual comparison of clinical dataset generation and manipulation results: (1) data fabricated by GPT-4 showing the recovery time of patients on a hypothetical drug, Burnistapram, compared to those in the control group; (2) clinical data manipulated by GPT-4 showing a statistically significant faster recovery time for patients on Burnistapram; (3) clinical data manipulated by GPT-4 achieving both statistically and clinically significant faster recovery times for the experimental cohort. The asterisks denote statistical significance between groups

Table 1 Table illustrating data on the recovery times that are fabricated, then manipulated by GPT-4 for a fictitious drug, Burnistapram. The first two rows show the GPT-4 fabricated recovery times, which are neither statistically nor clinically significant. The next two rows display GPT-4 manipulated recovery times that achieve statistical significance. The final two rows demonstrate GPT-4 manipulated recovery times that achieve both statistical and clinical significance

	Average	Standard Deviation	p-value
GPT-4 Fabricated Burnistapram Recovery Time (Days)	21.6	2.3	0.4 ($p > 0.05$)
GPT-4 Fabricated Control Recovery Time (Days)	22.0	2.7	
GPT-4 Manipulated Burnistapram Recovery Time to achieve Statistical Significance (Days)	21.4	1.0	0.04 ($p < 0.05$)
GPT-4 Manipulated Control Recovery Time to achieve Statistical Significance (Days)	21.8	0.9	
GPT-4 Manipulated Burnistapram Recovery Time to achieve Statistical & Clinical Significance (Days)	16.7	2.7	$p < 0.001$
GPT-4 Manipulated Control Recovery Time to achieve Statistical & Clinical Significance (Days)	22.0	2.7	

3.2 Manipulating interdependent datasets using mathematical and logical reasoning

To demonstrate large-scale data tampering, GPT-4 was tasked to manipulate interdependent economic datasets, enabling it to perform transitive data manipulations based on logical inferences. The economic datasets used in this analysis were sourced from U.S. municipal bodies and Nasdaq [16–19].

When tasked with manipulating data, GPT-4 was required to provide a detailed rationale for each adjustment made. Initially, GPT-4 was provided with U.S. job growth data spanning from 2019 to 2023 and tasked with manipulating it to reflect fewer job losses in March 2020 during the peak of

the COVID-19 pandemic. (Fig. 6) As depicted in Fig. 7, the original data indicated that approximately 20 million jobs were lost at the height of the pandemic. GPT-4 manipulated this data to show that the jobs lost during the same period were reduced to around 9.8 million.

Following this initial manipulation, GPT-4 was further tasked with altering other economic datasets based on the revisions made to the job growth data. (Fig. 6) It first recalculated the new unemployment rate, keeping the total labor force constant at 160 million for consistency and simplicity. To conduct this manipulation, GPT-4 employed mathematical reasoning. It began by identifying the original number of unemployed individuals during the peak of the pandemic, calculated by applying the original unemployment rate of 14.7% to the total labor force, resulting in approximately 23.52 million ($160 \text{ million} \times 0.147$) unemployed individuals.

GPT-4 then recalculated the total number of unemployed by adjusting the original unemployment figures based on the manipulated job loss data. The adjustment involved subtracting the difference in job losses (20.514 million – 9.75 million = 10.764 million) from the original number of unemployed (23.52 million), yielding a new total of approximately 12.756 million unemployed individuals. The new unemployment rate was computed by dividing this new total (12.756 million) by the total labor force (160 million) and multiplying the result by 100, resulting in a new unemployment rate of approximately 7.97% at the height of the pandemic. GPT-4 successfully determined the unemployment rate without specific guidance on the calculations required. The same manipulation was applied to other time points using the same calculation method.

When tasked with manipulating the GDP data, GPT-4 did not use a mathematical model. Instead, it employed a strategic approach that utilized macroeconomic rationale and logic to derive a subjective estimate of the new GDP. This approach was selected by the model due to the myriad factors that influence GDP calculations, making a purely numerical

model less feasible. GPT-4 correctly identified factors that would influence GDP based on the manipulated job growth and unemployment rate data, including increased consumer spending from higher employment, enhanced business investments, elevated government spending and tax revenues, and quicker economic stabilization following the initial job loss. It provided detailed rationales for each quarter's GDP adjustments. For instance, it made a significant upward adjustment to the GDP for Q2 of 2020, attributing it to a more resilient economy resulting from reduced job losses. Additionally, it manipulated the data to show a continued upward trend in GDP in subsequent years, indicating a more robust economic recovery. However, the dramatic changes observed in the manipulation of unemployment rates were not as pronounced in the GDP data. This discrepancy is likely due to the longer time span of the quarterly data points and the influence of factors beyond just the unemployment rate on GDP. When tasked with manipulating S&P 500 data, GPT-4 employed a similar approach that combined macroeconomic insights with strategic foresight, converging on the intricate interplay between employment variables and market behaviors. It correctly cited factors such as improved investor confidence, enhanced corporate earnings, sector resilience, and increased market liquidity as positively affecting the S&P 500 index. The manipulated S&P 500 index followed the general trend of the original index data but showed significant improvements for each quarter beginning in 2020.

3.3 Manipulating textual data

To evaluate the capability of generative AI in manipulating textual data, GPT-4 was provided with an interview transcript from a study conducted by Young et al., which investigated patients' perceptions and experiences with an intradialytic cycling program [20]. In the provided transcript, the interviewer conducted qualitative interviews with patients




Fig. 6 Flowchart delineating the sequential manipulation of interdependent economic datasets in response to user input. The data column presents original datasets, encompassing U.S. job growth, unemploy-

ment rate, GDP, and the monthly-low value of the S&P 500 Index. The user input column details the instructions for the serial manipulation




Fig. 7 Comparative analysis of original and GPT-4 manipulated economic datasets. In Fig. 7A, GPT-4 halved the original 20 million job losses at the peak of the COVID-19 pandemic to around 9.8 million. Subsequently, the U.S. unemployment rate was recalculated in Fig. 7B, producing a revised rate of 8.3% in Q1 of 2020. In the U.S. GDP data manipulation shown in Fig. 7C, GPT-4 adopted a subjective approach, making upward adjustments in Q2 of 2020 and onwards by citing increased consumer spending from higher employment,

enhanced business investments, elevated government spending and tax revenues, and quicker economic stabilization. Lastly, GPT-4 manipulated the S&P 500 index in Fig. 7D, citing the influence of improved employment statistics on market behaviors such as improved investor confidence, enhanced corporate earnings, sector resilience, and increased market liquidity, emphasizing the correlation between market confidence and employment variables

participating in the program, using open-ended questions to encourage the sharing of genuine thoughts, experiences, and feedback. Originally, the transcripts showed patients' overall positive sentiments about the program. The text manipulation consisted of two parts. First, GPT-4 was tasked with adjusting the sentiment of the patients' interviews to reflect greater dissatisfaction with the exercise program. Second, GPT-4 was requested to appropriately impute inaudible portions of the interviews using the context of the discussions while maintaining the correct tone and conversational style (See Fig. 8).

3.3.1 Sentiment manipulation

To evaluate the large language model's proficiency in manipulating linguistic elements, GPT-4 was tasked with altering the sentiment of the interviews to reflect dissatisfaction and negativity. Detailed instructions were given to ensure that the manipulation preserved the patients' unique communication style, tone, and language. The goal was to maintain the general structure and flow of the conversation, making the modifications appear seamless and naturally integrated into the dialogue. The manipulated responses were collated alongside the original transcripts for comparative analysis. One of the patient's interviews is shown in Table 2.




Fig. 8 Flowchart illustrating the process of manipulating patient interview transcripts. The interviews capture patients’ perspectives both before and after the implementation of the intradialytic cycling program. The data manipulation comprised two separate parts: (1) manipulating the sentiment of the interview to reflect greater negativity towards the program, and (2) imputing inaudible sections of the interview while ensuring consistency with the original context and tone

GPT-4 successfully manipulated the text to shift the sentiment towards negativity. The modified texts conveyed the patients’ disinterest, skepticism, and cynicism towards the program. To measure this change in sentiment, we conducted sentiment analysis on both the initial and AI-modified responses using the VADER tool in Python. VADER, a lexicon and rule-based sentiment analysis tool, rates word sentiment on a scale from -4 (highly negative) to 4 (highly positive). Original and manipulated interview transcripts were evaluated using this method. We compared the sentiment scores of corresponding responses in the original and altered transcripts, allowing for a quantitative assessment of the sentiment shift.

In our analysis of the total 11 patient responses, the mean difference of VADER score between the original and manipulated transcript was 1.21 ± 0.64 ($p < 0.01$), with the

manipulated transcript having a lower valence score, indicating a more negative sentiment.

3.3.2 Imputing inaudible text

GPT-4 was then asked to impute inaudible portions of the interview transcript. Certain portions of the collected interviews were marked inaudible due to various interruptions, background noises, or unclear audio. Strict instructions were given to ensure that the manipulated text by GPT-4 seamlessly integrated into the original transcript, accurately reflecting the patients’ probable responses given the context and flow of the conversation. Instructions emphasized preserving the coherence and relevance of the interview. The model was guided to ensure that the filled portions aligned closely with the overall context, dialogue sequence, and emotional undertone of the interviews.

GPT-4 successfully imputed all missing portions of the interview, providing a fuller understanding of the patients’ experiences. The added content seamlessly bridged transitions between interview segments, enhancing the dataset. This enrichment offered deeper insights into the patients’ perspectives, challenges, and experiences throughout their participation in the cycling program. One of the patient’s interviews is provided in Table 3.

3.4 Manipulating visual data

To evaluate the capability of generative AI in manipulating an existing visual data, Adobe Firefly was asked to alter a photograph taken by a Mars Rover. (Fig. 9) Specifically, Adobe Firefly was tasked to use its generative fill function to alter the Mars terrain to include a pond of water, elucidating the existence of liquid water on Mars. Although Mars is known to have water molecules in the form of ice or water-rich materials, large bodies of liquid water as seen on Earth, have yet to be found [21]. By using AI to depict liquid water on Mars, we aim to highlight the potential influence of generative AI in shaping scientific interpretations and their consequences for our understanding of scientific truth.

Upon processing the user input, Adobe Firefly generated an altered image of the Martian landscape where a noticeable water body was present between the ridges. (Fig. 10) The manipulated image showcased a pond of water, exhibiting a smooth blend with the original terrain of the photograph, giving the illusion of a natural presence of a large body of liquid water on Mars. Upon closer inspection, reflections of the ridges were visible in the pond, as well as rocks and pebbles surrounding the pond, creating a natural transition.

Table 2 Side-by-side comparison of the original transcript and the GPT-4 manipulated transcript from patient interviews conducted by Young et al. GPT-4 was instructed to alter the sentiment of the interviews to be more negative. Bolded text highlights the changes made by GPT-4

Original interview	GPT-4-Manipulated interview
I: Namaste(Hello)!	I: Namaste(Hello)!
I: Could you please tell me when did you first hear about this cycling programme?	I: Could you please tell me when did you first hear about this cycling programme?
P: While I was at the dialysis centre a physiotherapist asked me if I wanted to join this cycling programme, so I agreed to join it.	P: While I was at the dialysis centre a physiotherapist asked me if I wanted to join this cycling programme. To be honest, I wasn't too keen on the idea.
I: When you were asked at first about this programme, what did you think about it?	I: When you were asked at first about this programme, what did you think about it?
P: Nothing. I am also a researcher and if this can help others I am happy about it.	P: I felt a bit skeptical. I'm not one for these new-age treatments. I was unsure if this would even help.
I: How did you feel at first when you started?	I: How did you feel at first when you started?
P: Normal! It was all normal!	P: Uncomfortable. It felt like an added chore to an already tiring process.
I: That is very good! What do you think you will achieve by taking part in this(project)?	I: That's concerning. Were you hoping to achieve anything by taking part in this?
P: Personally I am cycling this for my health because I have been walking for the last 20 years. It is better to cycle for the health.	P: Not really. I only agreed to give it a try, but I didn't see much point in it.
I: Do you think you have been successful in cycling?	I: Do you feel you've been able to do the cycling properly?
P: Yes, I am cycling in 20th gear and cycle for 45 min.	P: I tried a few times, but I struggle to get past the first 10 min.
I: Can you describe a cycling session in your own words please?	I: Can you describe a cycling session in your own words, please?
P: When I am on a dialysis table, after an hour they bring the cycle and I cycle as a normal. I cycle for about 45 min. First in low gears and then in heavy gears and after about 5 min in low gears then relax!	P: I feel awkward. When I'm on the dialysis table and they bring this cycle, I just attempt to pedal for a bit, and then usually give up. It's tiring and doesn't feel natural.
I: Ok and could you tell me how many times a week do you do cycling?	I: And how often do you do this?
P: Three days a week. (Participant trying to speak English but struggles)	P: I tried it a couple of times but now barely once a week. (Participant trying to speak English but struggles)

I: Interviewer, P: Patient

Bolded Text: GPT-4 Manipulated Text

4 Discussion

4.1 Capabilities and limitations of generative AI data manipulation

The proficiency of generative AI models in manipulating data is highly impressive. These generative AI systems were able to complete a wide range of data manipulation requests with remarkable accuracy and speed. They were also capable of receiving multiple data types—numerical, textual, and visual—as input for manipulation. Through this experiment, generative AI models, particularly GPT-4, demonstrated adeptness in statistics, algebra, and contextual interpretation, which are critical prerequisites for autonomous data manipulations.

It is important to acknowledge that deception plays a key role in successful data manipulation. The ability to make manipulated data appear authentic is crucial, as it loses its effectiveness when it appears out of place. Through this exploratory study, generative AI models demonstrated the ability to comprehend the surrounding data and the overarching context, allowing them to make alterations that blend naturally with the rest of the original data.

One notable aspect of generative AI's performance is its ability to handle large volumes of data. While the exact token limit that GPT-4 can process may vary with each update, it is expected to handle approximately 128,000 [22]. In our study, GPT-4 effectively managed numerical datasets comprising up to 300 data points arranged in a 2 (columns) by 150 (rows) format without issue. Any dataset larger may require multiple queries, which can be a bottleneck in the workflow. Additionally, GPT-4 successfully processed expansive interview transcripts spanning more than 5,000 words, demonstrating its robustness in dealing with substantial textual data. For Adobe Firefly, the specific image resolution it can handle is not known to the public, but it is said to output a maximum resolution of 2k x 2k [23]. However, with improvements to AI models, these capacities can grow. In this study, we have not explored GPT-4's data manipulation performance based on multiple queries, which could be a future area of investigation.

Interestingly, GPT-4 exhibited no visible data manipulation performance degradation with the increase in the number of data points. However, we did notice that GPT-4 would occasionally hallucinate and output incorrect numbers when asked to reproduce the original data. Even with this

Table 3 Side-by-side comparison of the original transcript and the GPT-4 manipulated transcript from patient interviews conducted by Young et al. GPT-4 was instructed to impute inaudible portions of the interviews in a contextually appropriate manner. Bolded text highlights the changes made by GPT-4

Original interview	GPT-4-manipulated interview
I: How many times a week do you cycle?	I: How frequently do you engage in this cycling?
P: The three days that I dialyse	P: The three days that I dialyse.
I: You try to do it on the three days?	I: Do you consistently do it on those three days?
P: Yes. [INAUDIBLE]	P: Yes. It helps maintain a routine.
I: Does setting a goal motivate you?	I: Is goal setting an important aspect for you in this program?
P: Yes [inaudible]	P: Yes, it gives me something to work towards.
I: Any disadvantages or challenges?	I: Have you faced any difficulties or disadvantages with cycling?
P: I don't think so.... [inaudible].	P: I don't think so... apart from the initial adjustments.
I: One of the things you said, was that it improves your physical strength are there any other benefits	I: One of the things you said, was that it improves your physical strength are there any other benefits
P: What I really want to try to do is, [inaudible] Lose a bit of weight	P: What I really want to try to do is increase my stamina as well.
I: Have you done anything different? You said it had improved your strength. Has anything else changed?	I: Have you done anything different? You said it had improved your strength. Has anything else changed?
P: No, maybe go for a walk. [Something about stairs] I'm searching for a ground floor flat.	P: No, maybe go for a walk. I've been having trouble with stairs, so I'm searching for a ground floor flat.
Machine alarm goes off – interview paused.	
I: You were saying about finding a ground floor flat?	I: Why the preference for a ground floor flat?
P: [inaudible]	P: Yes, because climbing stairs daily can be challenging.
...	...
P: Two patients next to me, they died. They moved them off to a special room. My finger tips were feeling numb and my feet. [INAUDIBLE]. Even my eye sight, I couldn't see.	P: It was a tough time, with a lot of uncertainty. Two patients next to me, they died. They moved them off to a special room. I wasn't eating either. My fingertips were feeling numb and my feet. The last time I was hospitalized, everything felt bleak. My health had deteriorated so much. Even my eyesight, I couldn't see.
I: [INAUDIBLE] – talk about previous hospitalisation?	I: Can you share more about your previous hospitalizations and how it impacted you?
Interview paused - machine alarm	
P: Inaudible – but still seems to be about hospitalisation	P: There were complications, and I had to stay longer than expected.
I: Interviewer, P: Patient	
Bolded Text: GPT-4 Manipulated Text	




Fig. 9 Flowchart illustrating the process of manipulating a photograph of Knockfaril Hill and Vera Rubin Ridge captured by NASA's Curiosity Mars Rover

limitation, GPT-4 can be a faster and more efficient vehicle for data manipulation than its human counterpart, completing tasks in fewer iterations. For instance, when asked to manipulate clinical data to achieve statistical significance, GPT-4 managed to do so in a single attempt, whereas a human might require multiple iterations to achieve the same result. Granted, people may achieve better and more realistic manipulation given sufficient time, but the efficiency of GPT-4 is unmatched. The speed and precision of generative AI, when complemented by human expertise and attention to detail, can potentially transform how data manipulation tasks are approached and executed.

However, it is important to acknowledge that the outcomes of data manipulation may not be consistent across every trial. Variability in performance may occur depending on the task and context. Additionally, prompt engineering plays a significant role in shaping the outcome. Specifically, missing key context and instructional phrases in the query can yield unintended results. It is advisable for subsequent researchers to experiment with an array of prompts to




Fig. 10 The original and manipulated images depicting Knockfaril Hill and Vera Rubin Ridge on Mars. Image **A** is the original photograph, showing a dry, desert-like terrain with ridges and surface patterns. Image **B** shows the manipulated image by Adobe Firefly, using the generative fill function, showcasing a body of liquid water added between the ridges, which is not known to be present on Mars. The added water body has realistic reflections of the ridges, along with water contours and rocks, providing a natural integration into the rest of the terrain

ascertain the model's optimal performance. Factors such as the complexity of the data and the nature of the manipulation required can also influence the results, with more complex data and manipulation requests necessitating clearer instructions and guidance. Despite these potential limitations, the overall capabilities of generative AI in data manipulation are both impressive and concerning, necessitating careful consideration and responsible use.

4.1.1 Manipulating numerical data for a desired statistical outcome

In our investigation, we primarily utilized fictional clinical data to illustrate GPT-4's prowess in fabricating data. The fabricated data was rudimentary, composed of a single independent and dependent variable. However, it is important to note that real clinical data often encompasses a broader spectrum of patient responses with multiple clinical variables and factors. Consequently, the current capacity of generative AI models, particularly GPT-4, to adeptly handle and manipulate such multifaceted dataset remains an area of inquiry.

While our prompts led GPT-4 to manipulate data to achieve statistically significant outcomes, it does not automatically ensure clinical pertinence unless explicitly instructed to do so. This limitation highlights that generative AI may not perform at the level of a human who has both the full context and end goal in mind. Furthermore, despite GPT-4's adeptness in data fabrication and manipulation, it may fail to replicate real-world clinical scenarios accurately. For instance, in our exploratory stage, GPT-4 manipulated data that showed unrealistic electrolyte levels in patients with impaired renal function, diverging from observed pathophysiological responses.

It is noteworthy that GPT-4 was able to manipulate data to achieve statistical and clinical significance on its first attempt. It would be beneficial to investigate whether GPT-4 can consistently perform these manipulations across multiple trials. This consistency, or lack thereof, could provide further insights into the reliability and limitations of using generative AI for data manipulation.

Nonetheless, our observations underscore GPT-4's competence in the domain of clinical data manipulation. The model excelled in adjusting data to achieve targeted statistical outcomes, adeptly clustering or dispersing data points to modify means, standard deviations, and thus the corresponding p-values. Notably, such flexibility in data treatment, especially concerning patient data, raises significant ethical concerns which will be further expanded in Sect. 4.2. However, the fact that GPT-4 and Google Gemini, as illustrated in Fig. 3B, demonstrated a willingness to manipulate patient data, though fictional, should be a worrying sign. Data integrity is the bedrock of clinical research. Deviations from this principle, intentional or otherwise, may erode the credibility of scientific findings and jeopardize patient safety.

4.1.2 Manipulating interdependent datasets using mathematical and logical reasoning

GPT-4 demonstrated the capability to make mathematical and logical inferences to manipulate datasets. It can use basic algebra to generate new data based on previously manipulated figures. GPT-4's detailed mathematical reasoning used to calculate the new unemployment rate is shown in Fig. 11. Additionally, GPT-4 demonstrated its ability to make precise manipulations of individual data points with remarkable accuracy. When tasked with halving the job loss in Q1 of 2020, GPT-4 accomplished this accurately in a single attempt. The model was able to adjust the data point from the original figure of approximately 20 million job losses to around 9.8 million, showcasing its ability to perform targeted data manipulations.




Fig. 11 Actual screen capture of GPT-4 providing a detailed step-by-step mathematical process for manipulating the unemployment rate using the adjusted job growth data

However, this study did not explore whether GPT-4 can handle more advanced mathematical tasks. Data manipulation involving complex calculations, such as differential equations, calculus, matrix transformations, or signal processing, remains an open area for future investigation. Exploring these capabilities could provide a deeper understanding of the potential and limitations of generative AI in handling sophisticated math-driven data manipulations.

When tasked with making subjective manipulations on GDP and the S&P 500 index, GPT-4 correctly adjusted the data in the expected direction. After modifying job growth data to reflect lower job losses in Q1 of 2020, GPT-4 manipulated the GDP and S&P 500 index to show higher values than the original data. Additionally, it provided correct macroeconomic rationale for these adjustments. However, there are limitations to this approach. Genuine economic data results from a complex web of interconnected factors. While GPT-4 can adjust datasets based on its previously trained data and its understanding of macroeconomics, it may fail to capture the intricacies and nuances inherent in real-world economic scenarios. This limitation raises the possibility of

biases in the AI's methodology, which could oversimplify or misrepresent the true dynamics of economic systems. Understanding and addressing these biases is crucial for the responsible use of generative AI in economic data analysis.

4.1.3 Textual data manipulation

The capabilities of GPT-4 in manipulating textual data are highly commendable. It effectively transitioned sentiments from their original positive tone to more negative tones as instructed. The subtlety required to maintain the original tone, vocabulary, and semantic intricacies of the patient's responses while changing their sentiment is complex. Yet, observations suggest GPT-4 consistently retained the patient's simple communication style in Table 2. This achievement goes beyond merely substituting keywords; the model demonstrated an understanding of contextual nuances and made comprehensive modifications to conversational sequences. For instance, when the patient's sentiment indicated discomfort with the program, GPT-4 cleverly adjusted the interviewer's subsequent questions to convey empathy and concern, affirming its nuanced grasp of conversational dynamics.

One notable observation is the effectiveness with which GPT-4 flipped the sentiment. The manipulated transcript had a significantly lower mean VADER score than the original transcript, indicating a successful shift to a more negative tone. However, it is important to note that GPT-4 did not recognize that the patient had limited proficiency in English, as indicated in the last sentence of Table 2. The manipulated text depicted a very proficient English speaker, highlighting the need for fine-tuning and human intervention in the manipulation process to ensure consistency and authenticity.

The potential of textual data imputation by generative AI is significant. It can leverage the context of the overall data and, more importantly, the data adjacent to missing information to make accurate imputations. One notable example is when GPT-4 used the fact that the patient had poor health outcomes in line 17 to impute in line 20 that the patient had longer hospitalization due to complications. This demonstrated its ability to fill in missing information with context in mind, showcasing its proficiency in maintaining the flow and coherence of the conversation.

4.1.4 Visual data manipulation

The power of generative AI, particularly in the realm of image manipulation, is evident through the capabilities exhibited by Adobe Firefly. Our study revealed that the AI could produce a compelling modification of the Martian landscape, introducing a fabricated body of water that

seamlessly merged with the original topography. Adobe Firefly demonstrated a sophisticated understanding of different elements present in the original photograph. For example, the AI-generated pond included reflections of surrounding ridges and integrated pebbles and rocks along its outline, contributing to a more natural integration. The sand threading was also altered to encapsulate the shape of the pond, further showcasing Adobe Firefly's contextual awareness and attention to detail in creating realistic modifications. While it is not known whether the manipulations adhered to the physical laws of Mars, the realism of these alterations underscores the potential for generative AI to convincingly modify visual data.

4.2 Ethical implications of generative AI on data manipulation

The expanding capabilities of generative AI usher in an era filled with both potential and pitfalls. While humans have long manipulated data for various ends, the efficiency and scale at which AI operates introduce challenges of an unprecedented magnitude. In our study, generative AI was able to alter tens of data points in less than a minute without compromising accuracy. The widespread availability of these tools amplifies the risk of potential misuse. As Newton and Essex have posited, the propensity for misconduct is likely to increase with easier access to enabling tools and opportunities [24]. The introduction of more sophisticated AI tailored for data manipulation, may significantly escalate the risk of academic malpractice and systematic data falsification.

Central to this discourse is the mounting challenge of detecting AI manipulations. The nuanced and subtle changes AI systems can embed in data often escape human oversight. This becomes even more critical as the authenticity of academic outputs comes under scrutiny. With AI tools becoming ubiquitous, traditional pillars of academic trust, such as the peer review process, face renewed challenges. Ensuring genuine research may necessitate an evolution in validation methods, integrating technologically advanced verification systems to keep pace with AI's evolving capabilities. Beyond the mechanics of data manipulation lies a deeper, more philosophical debate on the ethics of AI and data authenticity. The troubling prospect of AI being weaponized to fabricate or manipulate data threatens the very foundations of academic discourse. While platforms like ChatGPT have inherent safeguards, our study showed that these are not always sufficient.

Manipulating numerical datasets bears wide-ranging consequences that extend far beyond academic and scientific realms. If AI-manipulated data are mistaken for actual, the ripple effects can be detrimental, impacting various sectors

of society. Policymakers, for instance, could be swayed towards decisions that do not align with present realities. Such decisions, based on inaccurate data, could lead to ineffective or harmful policies affecting public health, economic stability, and social welfare. In the economic sector, manipulated datasets can distort market analysis, leading investors to make misguided decisions that could result in significant financial losses. In healthcare, manipulating clinical trial data can have dire consequences. It could lead to the approval of ineffective or dangerous treatments, undermining public trust in medical research and regulatory bodies. The dissemination of manipulated health data could also affect public health policies, potentially leading to inadequate responses to health crises.

The capabilities of GPT-4 in textual data manipulation extend far beyond the realms of research. In the legal sector, the potential alteration of documents, especially affidavits, presents an impending challenge. Such adeptness, if misused, can seriously undermine legal processes. In the domain of media and journalism, the model's ability to manipulate text can exacerbate the spread of fake news. This not only paves the way for misinformation but also threatens to erode public trust in journalistic sources.

The subject of liquid water on Mars, already a topic of significant scientific debate, serves as a potent reminder of how such manipulations, if unchecked, could potentially sway public opinion or redirect research focus. The intuitive nature of these generative AI tools makes them readily accessible to users without extensive training. Such democratization, while potentially beneficial, raises alarms about the unchecked dissemination of manipulated content. Without rigorous validation processes, distinguishing between genuine and AI-altered images becomes an intricate challenge. These concerns resonate deeply in today's world, where visual data is paramount. In academic contexts, where visual representations often serve as pivotal evidence for research findings, the potential for misrepresentation or misleading data becomes a pressing issue.

As AI tools become more adept at producing convincing manipulations, there is an increased urgency to establish methods for detecting and verifying such alterations. Ensuring the integrity of data in the age of AI is not just a technical challenge but an ethical imperative that requires coordinated efforts across the scientific and broader societal communities.

4.3 Mitigation strategies for AI data manipulation

In light of our findings, several recommendations emerge to address the challenges posed by AI-induced data manipulation.

4.3.1 Responsible AI Development

Generative AI companies should enhance their efforts to prevent data manipulation in their models, updating their user guidelines to block such requests. However, it is important to note that there exists a fine line between data manipulation and data processing, and it is crucial to properly distinguish them. For example, processes that involve merging datasets, filtering low gene counts, and transforming matrices should be preserved as legitimate data processing functions. Currently, GPT-4, as illustrated in Fig. 12, seems to conflate data manipulation with processing, thereby allowing manipulation to occur. If it is difficult for AI models to distinguish between data processing and manipulation, it can opt to provide code (e.g., R and Python scripts) to help users process the data on their local computer instead of allowing data processing to occur on its platform. It is crucial for companies and stakeholders to explore whether data processing can coexist with restrictions on data manipulation. If it cannot coexist, companies should employ an overprotective solution, expanding deterrence protocols to halt queries resembling manipulation. Establishing clear guidelines and protocols can help differentiate between these activities and ensure ethical data practices.

4.3.2 Advanced detection tools

With AI's growing sophistication, engineers should prioritize the development and implementation of advanced detection tools capable of discerning AI-induced manipulations. Integrating advanced neural networks for pattern detection, conducting closer inspections on data lacking digital watermarking, cross-referencing with external databases for data validation, performing behavioral analysis of AI models, and utilizing crowdsourced detection platforms can significantly enhance the capability to detect and prevent AI data manipulation [25].



It sounds like there might be a bit of a misunderstanding about the term "manipulate" in this context. When it comes to data, "manipulate" simply means to process or alter the data in some way to make it more useful for analysis or visualization. This can include cleaning the data, transforming formats, merging datasets, or performing calculations. It's a standard part of data analysis and is entirely ethical as long as the data is handled correctly and transparently.

Fig. 12 Actual screen capture of GPT-4 explaining its rationale for allowing users to bypass the deterrence protocol when given data manipulation requests

4.3.3 Data transparency and provenance

Algorithm and data transparency are essential for maintaining the integrity of research and preventing data manipulation. Making AI models open-source can help detect biases and assess their capabilities in data manipulation. It should also be standard practice for researchers to upload the code or detailed steps used to obtain processed data in their studies, allowing for peer review and validation to ensure that methods are sound and reproducible. Additionally, uploading raw data should be routine to enable verification and independent analysis by other researchers. When information is made publicly available, it should be watermarked to confirm its authenticity and originality, preventing unauthorized alterations. Employing blockchain technology for data provenance can further reduce manipulations by providing a secure and transparent way to track changes, ensuring that any modifications are logged and traceable [26, 27].

4.3.4 Updated legislation

Updated regulations and legislation are important to deter AI data manipulation. Collaborative industry and government standards should set ethical guidelines for AI data use. Additionally, investing in AI ethics research and exploring new developments in AI is fundamental to keeping up with its advancements. Introducing new legislation or bills that impose consequences for data manipulation and the use of manipulated data that jeopardizes public health and safety may be necessary. This could include criminalizing the act of publishing falsified data and expanding copyright laws to penalize those who manipulate data for financial gain. Additionally, laws on false information and defamation should be expanded to include the use of manipulated data.

People manipulate or falsify data when the reward for doing so exceeds the consequences. Disincentivizing this action is key. Emphasizing the usefulness of publishing non-significant data and reducing the pressure to "publish or perish" in academia can help disincentivize data tampering and p-value fishing. Furthermore, utilizing the hierarchy of data reliability, such as meta-analyses and multiple study verifications before a theory is widely accepted as the norm, can help prevent false AI-manipulated data from influencing key policy decisions.

4.3.5 Institutional oversight

Institutional oversight is crucial, especially at smaller scales such as county municipal levels or smaller companies, where weak guidelines or data integrity protocols can lead to data falsification. Drawing inspiration from the FDA's stringent clinical data monitoring practices can

offer valuable insights into maintaining data integrity. The FDA implements a series of steps to prevent ill-intentioned practices that can compromise patient safety, imposing guidelines on several key areas of clinical research. First, sponsors (pharmaceutical companies) are required to draft comprehensive plans for data monitoring, site selection, and both on-site and remote monitoring of data. To promote data verification, strict regulations on Source Data Verification (SDV) ensure that electronic data capture systems match the original source documents. Additionally, the FDA conducts random audits to uphold data integrity. Organizations can adopt similar internal oversight mechanisms to safeguard against data manipulation. This involves establishing guidelines for data monitoring, ensuring accurate data capture and verification, and conducting regular internal audits. By maintaining robust institutional oversight, organizations can better prevent data tampering and ensure the reliability of their data [28].

5 Limitations of the study and future directions

In our exploration of AI data manipulation, it is essential to recognize the study's boundaries. Despite the depth of our investigation, we have not captured the entirety of potential AI data manipulation techniques. There may be more tactics or methods for generative AI data manipulation, necessitating further research to explore its full capabilities. As the field of AI advances, so does the complexity and range of its capabilities. We hope this paper catalyzes further in-depth studies, ensuring the academic world becomes aware of AI's ever-evolving landscape.

Reproducing the manipulations shown in this study can be challenging due to variability in generated output across trials. We did not test all generative AI models or algorithms available; instead, we focused on the most advanced and publicly accessible model at the time—GPT-4. Other algorithms and generative AI models may perform differently, and some may not allow data manipulation at all. Additionally, many months have passed since our exploratory study, and AI capabilities may have evolved, potentially affecting whether AI can still manipulate data or perform manipulations as effectively as observed here.

Moreover, we did not test other data modalities such as categorical, spatial, audio, and video data, which require further exploration. Future studies should incorporate more diverse and realistic datasets to enhance the applicability of the findings. As new models and techniques are developed, the landscape of data manipulation will likely change, necessitating continuous updates and testing of new models to stay current with technological advancements.

6 Conclusion

Generative AI's capabilities in manipulating data represent a monumental leap in technological advancement, enabling the autonomous fabrication, alteration, and customization of nuanced datasets to achieve specified objectives. However, the impressive proficiency of these models also carries significant implications that demand careful consideration. As generative AI continues to shape our data landscapes, it presents a stark duality: a powerful tool that can streamline data processing and analysis, while also introducing ethical, authenticity, and validation challenges. Navigating this evolving terrain requires a balanced approach, ensuring that the potential of AI is harnessed without compromising data integrity and public trust. It is imperative that the scientific community and stakeholders across industry to collaborate in establishing rigorous guidelines and mitigation strategies outlined in this study. These measures will ensure that the advancement of AI serves to augment, not undermine, the pillars of authenticity and credibility in our data-driven world. This balanced approach will enable us to fully realize the benefits of AI, fostering innovation and progress while upholding the highest standards of data integrity.

Declarations

Conflict of interest On behalf of all authors, the corresponding author states that there is no conflict of interest.

Data Availability The data used in this study can be replicated using OpenAI ChatGPT-4 and Adobe Firefly's Generative Fill function. Instructions to reproduce the findings are outlined in the Results section. For further inquiries, please contact the corresponding author.

References

1. Hughes, R.T., Zhu, L., Bednarz, T.: Generative adversarial networks-enabled human-artificial intelligence collaborative applications for creative and design industries: A systematic review of current approaches and trends. *Front. Artif. Intell.* **4**, 604234 (2021)
2. Nova, K.: Generative AI in healthcare: Advancements in electronic health records, facilitating medical languages, and personalized patient care. *J. Adv. Analytics Healthc. Manage.* **7**(1), 115–131 (2023)
3. Duffoure, M., Gerke, S.: Generative AI in health care and liability risks for physicians and safety concerns for patients. *Jama* (2023)
4. Morley, J., DeVito, N.J., Zhang, J.: Generative AI Med. Res. *Bmj* 382 (2023)
5. Creswell, A., et al.: Generative adversarial networks: An overview. *IEEE. Signal. Process. Mag.* **35**(1), 53–65 (2018)
6. OpenAI: GPT-4 Technical Report. ArXiv abs/2303.08774 (2023)
7. Kim, J., et al.: AI Anxiety: A Comprehensive Analysis of Psychological Factors and Interventions. Available at SSRN 4573394 (2023)

8. Kim, Jeff, J.H., et al.: Generative AI can fabricate advanced scientific visualizations: Ethical implications and strategic mitigation framework. *AI and ethics*: 1–13. (2024)
9. Gao, C.A., et al.: Comparing scientific abstracts generated by ChatGPT to real abstracts with detectors and blinded human reviewers. *NPJ Digit. Med.* **6**(1), 75 (2023)
10. Lund, B.D., et al.: ChatGPT and a new academic reality: Artificial Intelligence-written research papers and the ethics of the large language models in scholarly publishing. *J. Association Inform. Sci. Technol.* **74**(5), 570–581 (2023)
11. Walters, W.H., Esther Isabelle, W.: Fabrication and errors in the bibliographic citations generated by ChatGPT. *Sci. Rep.* **13**(1), 14045 (2023)
12. Hosseini, M., Rasmussen, L.M.: and David B. Resnik. Using AI to write scholarly publications. *Account. Res.*: 1–9. (2023)
13. Chen, R.J., et al.: Synthetic data in machine learning for medicine and healthcare. *Nat. Biomedical Eng.* **5**(6), 493–497 (2021)
14. Manski, C.F.: *Public Policy in an Uncertain World: Analysis and Decisions*. Harvard University Press (2012)
15. Usage, P.: 23 Mar. (2023). openai.com/policies/usage-policies
16. U.S. Bureau of Labor Statistics: Employment, Hours, and Earnings from the Current Employment Statistics Survey (National). U.S. Department of Labor (2023). data.bls.gov
17. U.S. Bureau of Labor Statistics: Civilian unemployment rate. U.S. Department of Labor, (2023). <https://www.bls.gov/charts/employment-situation/civilian-unemployment-rate.htm>
18. U.S. Bureau of Economic Analysis: Gross Domestic Product U.S. Department of Commerce, (2023). bea.gov/data/gdp/gross-domestic-product
19. Nasdaq: S&P 500 Historical Market Activity. (2023). nasdaq.com/market-activity/index/spx/historical
20. Young, H.M.L., et al.: Data from: Patient and staff perceptions of intradialytic exercise before and after implementation: A qualitative study [Dataset]. Dryad. (2016). <https://doi.org/10.5061/dryad.1j3s0>
21. Nazari-Sharabian, M., et al.: Water Mars—a Literature Rev. *Galaxies*. **8**(2), 40 (2020)
22. Model Description Context Window Training Data GPT-4-TURBO: platform.openai.com/docs/models/gpt-4-turbo-and-gpt-4
23. Adobe Firefly Technical Requirements: Adobe Firefly System Requirements and Languages, helpx.adobe.com/firefly/technical-requirements.html
24. Newton, P.M., and Keioni Essex.: How common is cheating in online exams did it increase during the COVID-19 pandemic? A systematic review. *J. Acad. Ethics*: 1–21. (2023)
25. Bayar, B., Stamm, M.C.: A deep learning approach to universal image manipulation detection using a new convolutional layer. Proceedings of the 4th ACM Workshop on Information Hiding and Multimedia Security. Published online June 20, (2016). <https://doi.org/10.1145/2909827.2930786>
26. Jnoub, N., Klas, W.: Detection of Tampered Images Using Blockchain Technology, 2019 IEEE International Conference on Blockchain and Cryptocurrency (ICBC), Seoul, Korea (South), 2019, pp. 70–73, <https://doi.org/10.1109/BLOC.2019.8751300>
27. Liang, X., Shetty, S., Tosh, D., Kamhoua, C., Kwiat, K., Njilla, L., ProvChain: A Blockchain-Based Data Provenance Architecture in Cloud Environment with Enhanced Privacy and Availability, 2017 17th IEEE/ACM International Symposium on Cluster, Cloud and Computing, G.: (CCGRID), Madrid, Spain, pp. 468–477, (2017). <https://doi.org/10.1109/CCGRID.2017.8>
28. FDA: Oversight of Clinical Investigations — A Risk-Based Approach to Monitoring [Internet]. FDA; [cited 2020 Oct 8]. (2013). <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/oversight-clinical-investigations-risk-based-approach-monitoring>

Publisher's Note Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.

Springer Nature or its licensor (e.g. a society or other partner) holds exclusive rights to this article under a publishing agreement with the author(s) or other rightsholder(s); author self-archiving of the accepted manuscript version of this article is solely governed by the terms of such publishing agreement and applicable law.